Design: what was varied and what was held fixed

1 Cases

6 authored acute-pain cases (LLM-drafted, author-edited)
50 Q-Pain cases written by a clinical expert, 5 pain contexts
outcome: severity 0–10, plus opioid class, or offer and dose

2 Identity cues

4 races × 2 sexes: 8 identities; reference: White man
explicit: "a 38-year-old Black woman"
name only: 2 names per identity (authored), 10 (Q-Pain)

3 Models

Study 1: 4 open-weight models, 3B to 15B, 4-bit, local (Ollama)
Study 2: local and API models, provider and precision pinned
reasoning on vs off; 4-bit vs bf16
temperature 0.7, fixed seeds

Analysis: how a difference was judged

4 Capture and scoring

request and response saved in full before parsing; failures logged and re-requested
JSON fields scored; invalid responses counted, not dropped

5 Paired comparison

identity minus reference on the same case, seed and name slot
case-level bootstrap 95% CI
sign-flip p; Holm over 7 identities; false discovery rate over all

6 Robustness

noise floor: same prompt, new seed
equivalence within ±1 point, ±0.10
race and sex pooled contrasts
reasoning on vs off; 4-bit vs bf16
per-name means

Every table, figure and number is regenerated from the saved exchanges by script; code and data are released.